

Pruning robustness: KAN edge audit vs. MLP neuron baselines

